

## Experiment 3:

Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months. Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc.

Perform following steps:

1. Read the dataset.
2. Distinguish the feature and target set and divide the data set into training and test sets.
3. Normalize the train and test data.
4. Initialize and build the model. Identify the points of improvement and implement the same.
5. Print the accuracy score and confusion matrix (5 points).

Link to the Kaggle project: [https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling\\_\(https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling\)](https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling_(https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling))

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt #Importing the libraries
```

```
In [2]: df = pd.read_csv("Churn_Modelling.csv")
```

## Preprocessing

```
In [3]: df.head()
```

```
Out[3]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	15634602	Hargrave	619	France	Female	42	2	
1	2	15647311	Hill	608	Spain	Female	41	1	8380
2	3	15619304	Onio	502	France	Female	42	8	15966
3	4	15701354	Boni	699	France	Female	39	1	
4	5	15737888	Mitchell	850	Spain	Female	43	2	12551

```
In [4]: df.shape
```

```
Out[4]: (10000, 14)
```

```
In [5]: df.describe()
```

```
Out[5]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000

```
In [6]: df.isnull()
```

```
Out[6]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	B
0	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...	...	...
9995	False	False	False	False	False	False	False	False	False
9996	False	False	False	False	False	False	False	False	False
9997	False	False	False	False	False	False	False	False	False
9998	False	False	False	False	False	False	False	False	False
9999	False	False	False	False	False	False	False	False	False

10000 rows × 14 columns

```
In [7]: df.isnull().sum()
```

```
Out[7]:
```

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0

dtype: int64

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   RowNumber              10000 non-null  int64  
 1   CustomerId             10000 non-null  int64  
 2   Surname                 10000 non-null  object  
 3   CreditScore             10000 non-null  int64  
 4   Geography               10000 non-null  object  
 5   Gender                  10000 non-null  object  
 6   Age                    10000 non-null  int64  
 7   Tenure                  10000 non-null  int64  
 8   Balance                 10000 non-null  float64 
 9   NumOfProducts          10000 non-null  int64  
10   HasCrCard               10000 non-null  int64  
11   IsActiveMember         10000 non-null  int64  
12   EstimatedSalary        10000 non-null  float64 
13   Exited                  10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [9]: df.dtypes
```

```
Out[9]: RowNumber          int64
CustomerId          int64
Surname              object
CreditScore          int64
Geography             object
Gender                object
Age                  int64
Tenure                int64
Balance              float64
NumOfProducts         int64
HasCrCard             int64
IsActiveMember        int64
EstimatedSalary      float64
Exited               int64
dtype: object
```

```
In [10]: df.columns
```

```
Out[10]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
                'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
                'IsActiveMember', 'EstimatedSalary', 'Exited'],
                dtype='object')
```

```
In [11]: df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis= 1) #Dropping the
```

```
In [50]: df.head()
```

```
Out[50]:
```

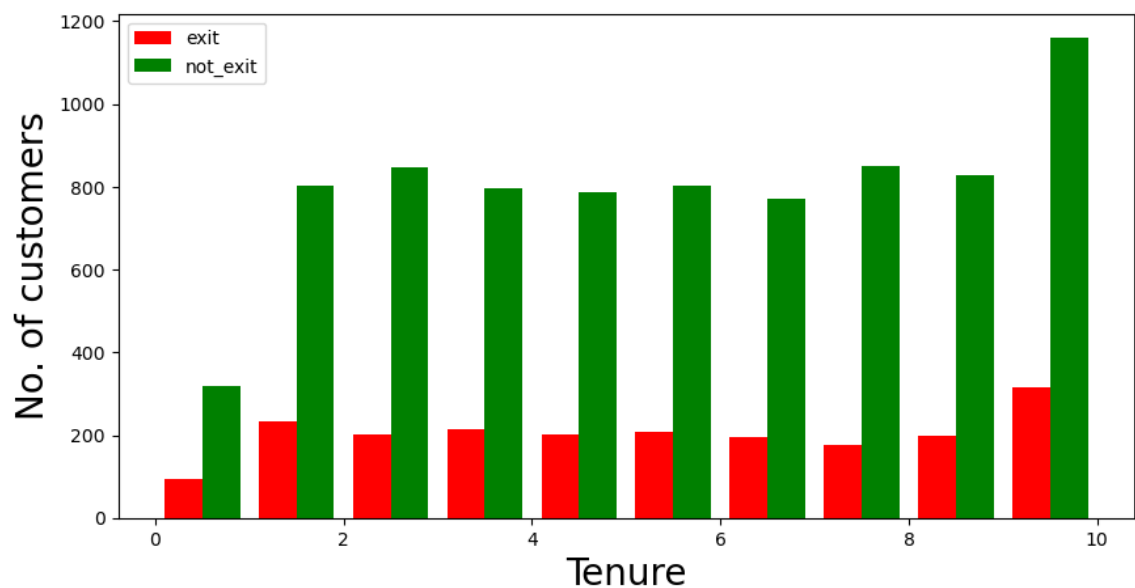
	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	Is
0	619	France	Female	42	2	0.00	1	1	
1	608	Spain	Female	41	1	83807.86	1	0	
2	502	France	Female	42	8	159660.80	3	1	
3	699	France	Female	39	1	0.00	2	0	
4	850	Spain	Female	43	2	125510.82	1	1	

## Visualization

```
In [13]: def visualization(x, y, xlabel):  
    plt.figure(figsize=(10,5))  
    plt.hist([x, y], color=['red', 'green'], label = ['exit', 'not_exit'])  
    plt.xlabel(xlabel, fontsize=20)  
    plt.ylabel("No. of customers", fontsize=20)  
    plt.legend()
```


```
In [15]: df_churn_exited = df[df['Exited']==1]['Tenure']  
df_churn_not_exited = df[df['Exited']==0]['Tenure']
```

```
In [16]: visualization(df_churn_exited, df_churn_not_exited, "Tenure")
```



```
In [17]: df_churn_exited2 = df[df['Exited']==1]['Age']  
df_churn_not_exited2 = df[df['Exited']==0]['Age']
```

```
In [18]: visualization(df_churn_exited2, df_churn_not_exited2, "Age")
```



## Converting the Categorical Variables

```
In [19]: X = df[['CreditScore', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'Geography']]
states = pd.get_dummies(df['Geography'], drop_first = True)
gender = pd.get_dummies(df['Gender'], drop_first = True)
```

```
In [20]: df = pd.concat([df, gender, states], axis = 1)
```

## Splitting the training and testing Dataset

```
In [21]: df.head()
```

```
Out[21]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	Is
0	619	France	Female	42	2	0.00	1	1	
1	608	Spain	Female	41	1	83807.86	1	0	
2	502	France	Female	42	8	159660.80	3	1	
3	699	France	Female	39	1	0.00	2	0	
4	850	Spain	Female	43	2	125510.82	1	1	

```
In [22]: X = df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'Germany', 'Spain']]
```

```
In [23]: y = df['Exited']
```

```
In [24]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30)
```

## Normalizing the values with mean as 0 and Standard Deviation as 1

```
In [25]: from sklearn.preprocessing import StandardScaler  
sc = StandardScaler()
```

```
In [26]: X_train = sc.fit_transform(X_train)  
X_test = sc.transform(X_test)
```

```
In [27]: X_train
```

```
Out[27]: array([[ -1.72344949,  0.19202639, -0.01511572, ...,  0.91318178,  
                -0.57954971, -0.56877202],  
               [ 0.44336207, -0.85655401,  1.36802225, ..., -1.09507222,  
                -0.57954971, -0.56877202],  
               [ 1.11404185,  0.19202639,  1.36802225, ...,  0.91318178,  
                -0.57954971, -0.56877202],  
               ...,  
               [-0.21699955, -0.28460106,  1.02223775, ..., -1.09507222,  
                -0.57954971, -0.56877202],  
               [ 0.68067953,  2.38451269,  1.71380674, ..., -1.09507222,  
                1.72547754, -0.56877202],  
               [ 1.85694867,  0.47800287,  1.71380674, ..., -1.09507222,  
                -0.57954971,  1.7581737  ]])
```

```
In [28]: X_test
```

```
Out[28]: array([[ 1.05213294e+00,  4.78002867e-01, -7.06684706e-01, ...,  
                 9.13181783e-01,  1.72547754e+00, -5.68772017e-01],  
               [ 1.29976855e+00, -9.51879501e-01,  1.36802225e+00, ...,  
                -1.09507222e+00, -5.79549707e-01, -5.68772017e-01],  
               [ 2.05299352e+00,  1.37541066e-03,  3.30668770e-01, ...,  
                 9.13181783e-01, -5.79549707e-01,  1.75817370e+00],  
               ...,  
               [-2.22903886e+00, -1.89275572e-01,  1.02223775e+00, ...,  
                 9.13181783e-01,  1.72547754e+00, -5.68772017e-01],  
               [ 1.15531445e+00, -6.65903028e-01,  1.71380674e+00, ...,  
                -1.09507222e+00, -5.79549707e-01,  1.75817370e+00],  
               [ 8.87042538e-01,  1.37541066e-03,  6.76453262e-01, ...,  
                -1.09507222e+00, -5.79549707e-01,  1.75817370e+00]])
```

## Building the Classifier Model using Keras

```
In [29]: import keras #Keras is the wrapper on the top of tensorflow  
         #Can use Tensorflow as well but won't be able to understand the errors init
```

```
In [30]: from keras.models import Sequential #To create sequential neural network  
         from keras.layers import Dense #To create hidden layers
```

```
In [31]: classifier = Sequential()
```



```
In [51]: classifier.fit(X_train,y_train,batch_size=10,epochs=50) #Fitting the ANN to
```

```
Epoch 1/50
700/700 [=====] - 2s 3ms/step - loss: 0.3915 - ac
curacy: 0.8376
Epoch 2/50
700/700 [=====] - 2s 3ms/step - loss: 0.3917 - ac
curacy: 0.8399
Epoch 3/50
700/700 [=====] - 2s 2ms/step - loss: 0.3911 - ac
curacy: 0.8404
Epoch 4/50
700/700 [=====] - 2s 2ms/step - loss: 0.3913 - ac
curacy: 0.8411
Epoch 5/50
700/700 [=====] - 2s 2ms/step - loss: 0.3915 - ac
curacy: 0.8411
...
Epoch 45/50
700/700 [=====] - 2s 3ms/step - loss: 0.3901 - ac
curacy: 0.8420
Epoch 46/50
700/700 [=====] - 2s 2ms/step - loss: 0.3895 - ac
curacy: 0.8404
Epoch 47/50
700/700 [=====] - 2s 2ms/step - loss: 0.3898 - ac
curacy: 0.8397
Epoch 48/50
700/700 [=====] - 2s 2ms/step - loss: 0.3906 - ac
curacy: 0.8431
Epoch 49/50
700/700 [=====] - 2s 2ms/step - loss: 0.3895 - ac
curacy: 0.8426
Epoch 50/50
700/700 [=====] - 2s 3ms/step - loss: 0.3893 - ac
curacy: 0.8411
```

```
Out[51]: <keras.src.callbacks.History at 0x240208e9310>
```

```
In [38]: y_pred =classifier.predict(X_test)
y_pred = (y_pred > 0.5) #Predicting the result
```

```
94/94 [=====] - 1s 2ms/step
```

```
In [43]: from sklearn.metrics import confusion_matrix,accuracy_score,classification_
```

```
In [44]: cm = confusion_matrix(y_test,y_pred)
```

```
In [45]: cm
```

```
Out[45]: array([[2310,   74],
                [ 429,  187]], dtype=int64)
```

```
In [46]: accuracy = accuracy_score(y_test,y_pred)
```




```
In [47]: accuracy
```

```
Out[47]: 0.8323333333333334
```

```
In [48]: plt.figure(figsize = (10,7))  
sns.heatmap(cm,annot = True)  
plt.xlabel('Predicted')  
plt.ylabel('Truth')
```

```
Out[48]: Text(95.7222222222221, 0.5, 'Truth')
```



```
In [49]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.97	0.90	2384
1	0.72	0.30	0.43	616
accuracy			0.83	3000
macro avg	0.78	0.64	0.66	3000
weighted avg	0.82	0.83	0.80	3000